

In-memory & Semi- External Memory Large- Scale Clustering

Disa Mhembere



K-means Motivation

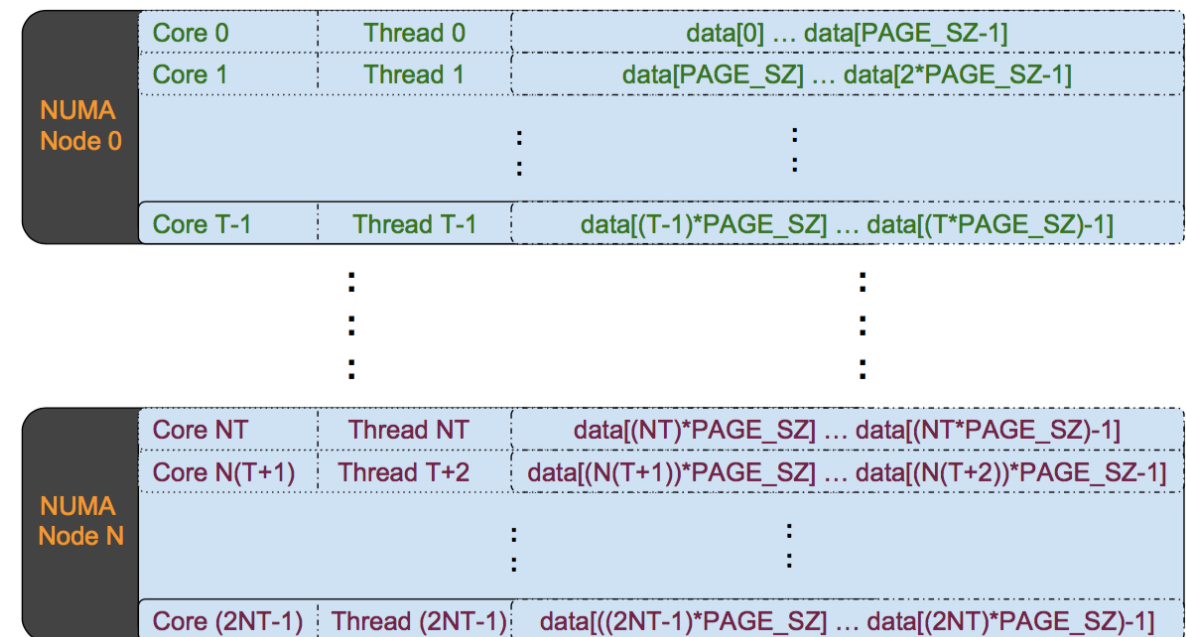
- One of the most widely used & important clustering techniques.
- Problems with k-means:
 - Historically does not scale well computation complexity: $O(kdn)$, Space complexity: $O(nd + kd)$
 - Others did not fully utilize CPU resources (toward computation efficiently)
 - Others have not take advantage of NUMA architecture nor new vectorization enabled CPUs

- Our approach:

- Scale-up not scale-out

- **The how** — Packages:

- k||means : Shared memory
- SEM-kmeans, Min-Triangle-SEM-kmeans: Semi-external memory



k||means

- Rethink Lloyd's EM steps. Why not (mostly) merge the two EM-steps into a super-step?
 - How? Share **no** data!
 - Per-thread data structures; combined recursively in || at end of super-step.
 - Result: Embarrassing ||ism
- Developed ||ized initializations via || kmeans++ or kmeans|| — they matter

Old memory bound: $O(nd + kd)$

New memory bound: $O(nd + Tkd)$, where $T = \#threads$

Algorithm 1 k||means algorithm

```
1: procedure K||MEANS( $V, C, K$ )
2:    $ptCentroids$  ▷ Per-thread centroids
3:    $clusterAssignment$  ▷ Shared, no conflict
4:   parfor  $\vec{v}_i \in V$  do
5:     for  $\vec{c}_i \in C$  do
6:        $[dist_{min}, cid_{min}] = \min(\mathbf{d}(\vec{v}_i, \vec{c}_i))$ 
7:     end for
8:      $ptCentroids[CURR\_THREAD][cid_{min}] += \vec{v}_i$ 
9:   end parfor
10:   $clusterMeans = \text{mergePtStructs}(ptClusters)$ 
11: end procedure

12: procedure MERGEPTSTRUCTS( $vectors$ )
13:   while  $|vectors| > 1$  do
14:      $PAR\_MERGE(vectors)$  ▷  $O(T \log n)$ 
15:   end while
16:   return  $vectors[0]$ 
17: end procedure
```



k||means Performance on Friendster eigs (66 Mil x 8)

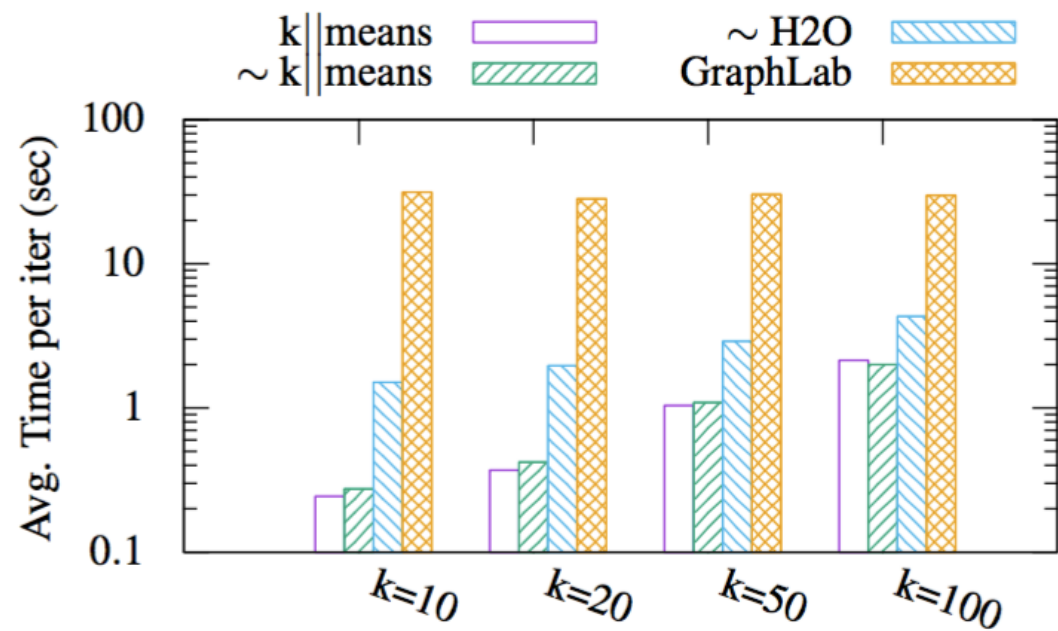


Fig 1: Log scale average time per iteration

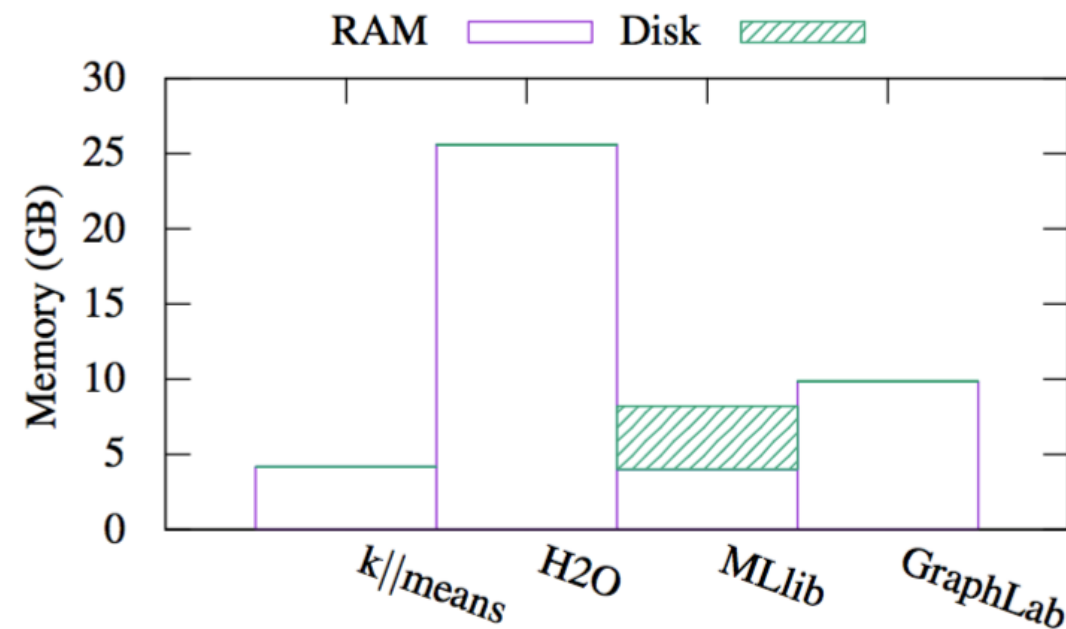


Fig 2: Log scale average time per iteration

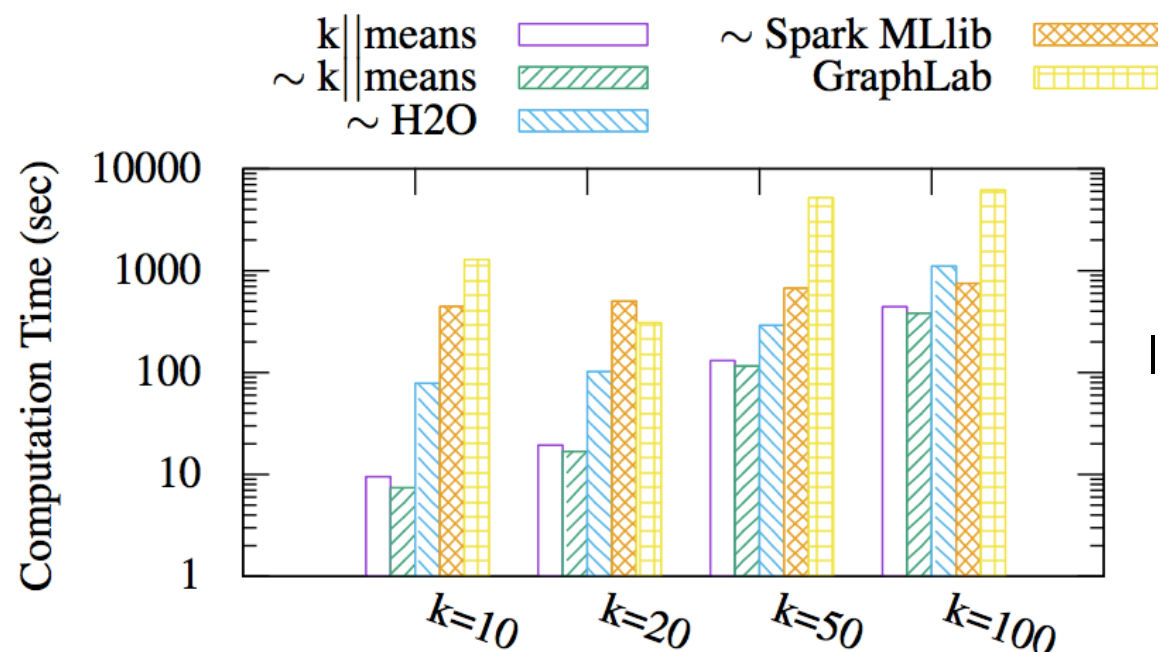


Fig 3: Log scale average computation time. Important because we see **greater** improvement than the per cost per iteration (Fig. 1) due to || initialization modules





SEM-kmeans

- Can we be comparable with less resources?
 - Semi-external:
 - Memory now $O(n + Tkd) < O(nd + kd) < O(nd + Tkd)$
 - $T = \text{\#threads/processes}$ and $Tkd \ll O(nd)$
- Further improve speed using a *modified* triangle inequality pruning algorithm with same memory bound: **Min-Triangle-SEM-Kmeans**
 - Performance improvement from 2 factors:
 - Reduction in # of distance computations
 - Reduction in I/O complexity



Min-Triangle-SEM-kmeans on Friendster

Micro Data: 66Mil x 8, Dense, Size: 4 GB, K=10



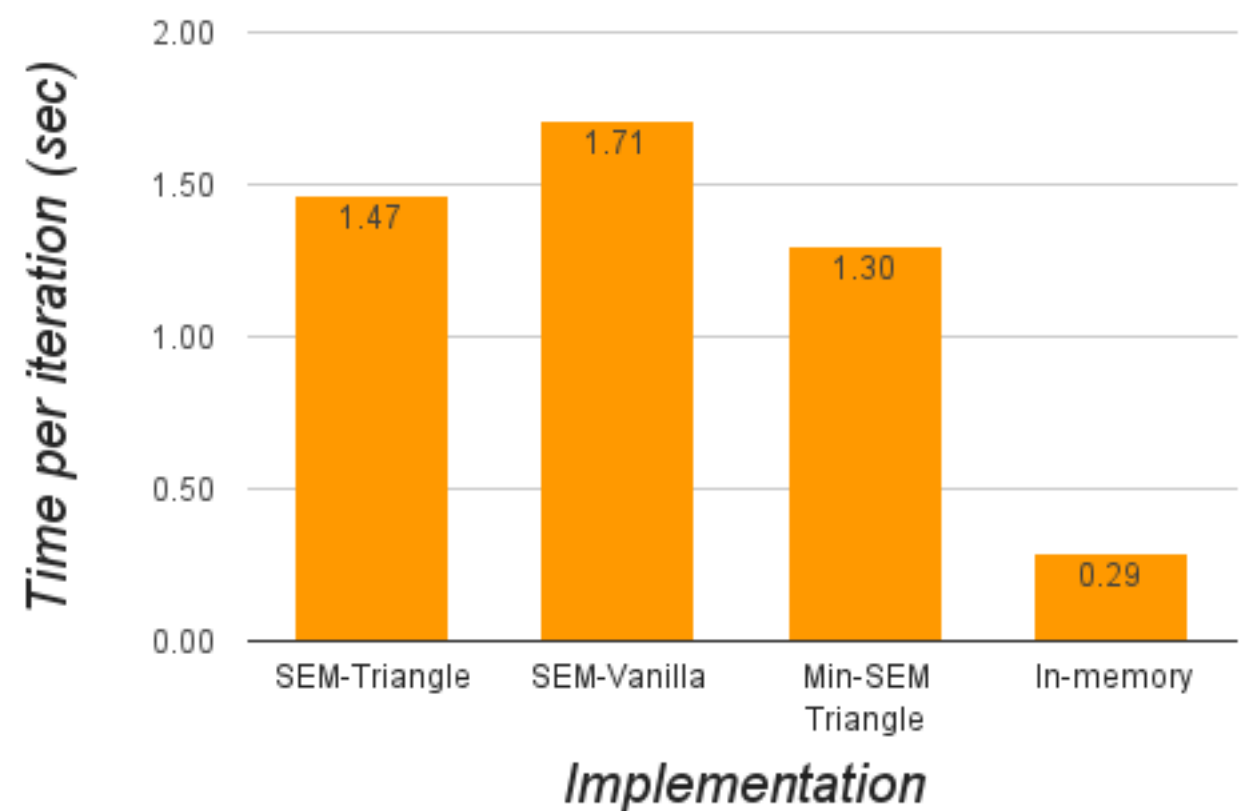

Insight: The effect of upper bound matrix, $O(nk)$ is minimal. It contributes only $< 2\%$ (Lemma 4 & Lemma 5)

Both depend a lot on the data ...
But still very I/O bound

Only **40% CPU** utilization ...



Variant comparison on Friendster



*More optimizations to follow,
1GB cache for all SEM



Applications

- Together with the FlashEigen-solver we can perform extremely fast and scalable spectral embedding!
 - Applies to:
 - Connectomics
 - Anomaly detection
 - Machine-Learning
 - NLP
 - A host of other interesting problems:
 - <https://sites.google.com/site/dataclusteringalgorithms/clustering-algorithm-applications>

